

# Recommender systems

## Making recommendations

### Predicting movie ratings

User rates movies using one to five stars

Ratings				
★				
★	★			
★	★	★		
★	★	★	★	
★	★	★	★	★

Movie	Alice(1)	Bob(2)	Carol(3)	Dave(4)
Love at last	5	5	0	0
Romance forever	5	?	?	0
Cute puppies of love	?	4	0	?
Nonstop car chases	0	0	5	4
Swords vs. karate	0	0	5	?

$$n_u = 4 \quad r(1,1) = 5$$

$$n_m = 5 \quad r(3,1) = 0 \quad y^{(3,2)} = 4$$

$n_u$  = no. of users

$n_m$  = no. of movies

$r(i,j) = 1$  if user  $j$  has rated movie  $i$

$y^{(i,j)}$  = rating given by user  $j$  to movie  $i$  (defined only if  $r(i,j) = 1$ )

## Using per-item features

### What if we have features of the movies?

$$n_u = 4$$

$$n_m = 5$$

$$n = 2$$

Movie	Alice(1)	Bob(2)	Carol(3)	Dave(4)	$x_1$ (romance)	$x_2$ (action)
Love at last	5	5	0	0	0.9	0
Romance forever	5	?	?	0	1.0	0.01
Cute puppies of love	?	4	0	?	0.99	0
Nonstop car chases	0	0	5	4	0.1	1.0
Swords vs. karate	0	0	5	?	0	0.9

$$x^{(1)} = \begin{bmatrix} 0.9 \\ 0 \end{bmatrix}$$

$$x^{(3)} = \begin{bmatrix} 0.99 \\ 0 \end{bmatrix}$$

For user 1: Predict rating for movie  $i$  as:  $w^{(1)} \cdot x^{(i)} + b^{(1)}$  ← just linear regression

$$w^{(1)} = \begin{bmatrix} 5 \\ 0 \end{bmatrix} \quad b^{(1)} = 0 \quad x^{(3)} = \begin{bmatrix} 0.9 \\ 0 \end{bmatrix}$$

$$w^{(1)} \cdot x^{(3)} + b^{(1)} = 4.95$$

→ For user  $j$ : Predict user  $j$ 's rating for movie  $i$  as  $w^{(j)} \cdot x^{(i)} + b^{(j)}$

## Cost function

Notation:

- $r(i,j) = 1$  if user  $j$  has rated movie  $i$  (0 otherwise)
- $y^{(i,j)}$  = rating given by user  $j$  on movie  $i$  (if defined)
- $w^{(j)}, b^{(j)}$  = parameters for user  $j$
- $x^{(i)}$  = feature vector for movie  $i$

For user  $j$  and movie  $i$ , predict rating:  $w^{(j)} \cdot x^{(i)} + b^{(j)}$

→  $m^{(j)}$  = no. of movies rated by user  $j$

To learn  $w^{(j)}, b^{(j)}$

$$\min_{w^{(j)}, b^{(j)}} J(w^{(j)}, b^{(j)}) = \frac{1}{2m^{(j)}} \sum_{i:r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{\lambda}{2m^{(j)}} \sum_{k=1}^n (w_k^{(j)})^2$$

number of features

# Cost function

牛啊

To learn parameters  $w^{(j)}, b^{(j)}$  for user  $j$  :

$$J(w^{(j)}, b^{(j)}) = \frac{1}{2} \sum_{i:r(i,j)} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{k=1}^n (w_k^{(j)})^2$$

To learn parameters  $w^{(1)}, b^{(1)}, w^{(2)}, b^{(2)}, \dots, w^{(n_u)}, b^{(n_u)}$  for all users :

$$J(w^{(1)}, \dots, w^{(n_u)}, b^{(1)}, \dots, b^{(n_u)}) = \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (w_k^{(j)})^2$$

合作

## Collaborative filtering algorithm → 通过协作预测电影因子

### Problem motivation

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)	$x_1$ (romance)	$x_2$ (action)
Love at last	5	5	0	0	?	?
Romance forever	5	?	?	0	?	?
Cute puppies of love	?	4	0	?	?	?
Nonstop car chases	0	0	5	4	?	?
Swords vs. karate	0	0	5	?	?	?

多用户协作  
预测新用户对电影的 W&B

$w^{(1)} = \begin{bmatrix} 5 \\ 0 \end{bmatrix}, w^{(2)} = \begin{bmatrix} 5 \\ 0 \end{bmatrix}, w^{(3)} = \begin{bmatrix} 0 \\ 5 \end{bmatrix}, w^{(4)} = \begin{bmatrix} 0 \\ 5 \end{bmatrix}$   
 $b^{(1)} = 0, b^{(2)} = 0, b^{(3)} = 0, b^{(4)} = 0$

using  $w^{(j)} \cdot x^{(i)} + b^{(j)}$

$w^{(1)} \cdot x^{(1)} \approx 5$   
 $w^{(2)} \cdot x^{(1)} \approx 5$   
 $w^{(3)} \cdot x^{(1)} \approx 0$

→  $x^{(1)} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$

是因为你有来自

### Cost function

Given  $w^{(1)}, b^{(1)}, w^{(2)}, b^{(2)}, \dots, w^{(n_u)}, b^{(n_u)}$

to learn  $x^{(i)}$  :

$$J(x^{(i)}) = \frac{1}{2} \sum_{j:r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{k=1}^n (x_k^{(i)})^2$$

→ To learn  $x^{(1)}, x^{(2)}, \dots, x^{(n_m)}$  :

$$J(x^{(1)}, x^{(2)}, \dots, x^{(n_m)}) = \frac{1}{2} \sum_{i=1}^{n_m} \sum_{j:r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2$$

# Collaborative filtering

Cost function to learn  $w^{(1)}, b^{(1)}, \dots, w^{(n_u)}, b^{(n_u)}$ :

$$\min_{w^{(1)}, b^{(1)}, \dots, w^{(n_u)}, b^{(n_u)}} \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (w_k^{(j)})^2$$

Cost function to learn  $x^{(1)}, \dots, x^{(n_m)}$ :

$$\min_{x^{(1)}, \dots, x^{(n_m)}} \frac{1}{2} \sum_{i=1}^{n_m} \sum_{j:r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2$$

Put them together:

$$\min_{\substack{w^{(1)}, \dots, w^{(n_u)} \\ b^{(1)}, \dots, b^{(n_u)} \\ x^{(1)}, \dots, x^{(n_m)}}} J(w, b, x) = \frac{1}{2} \sum_{(i,j):r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (w_k^{(j)})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2$$

再次整合

## Gradient Descent

collaborative filtering

Linear regression (course 1)

repeat {

$$w_i = w_i - \alpha \frac{\partial}{\partial w_i} J(w, b, x)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b, x)$$

$$w_i^{(j)} = w_i^{(j)} - \alpha \frac{\partial}{\partial w_i^{(j)}} J(w, b, x)$$

$$b^{(j)} = b^{(j)} - \alpha \frac{\partial}{\partial b^{(j)}} J(w, b, x)$$

$$x_k^{(i)} = x_k^{(i)} - \alpha \frac{\partial}{\partial x_k^{(i)}} J(w, b, x)$$

}

parameters  $w, b, x$   $x$  is also a parameter

Binary labels: saves, likes and clicks

Binary labels

将线性  
转为  
二分类

Movie	Alice(1)	Bob(2)	Carol(3)	Dave(4)
Love at last	1	1	0	0
Romance forever	1	? ←	? ←	0
Cute puppies of love	? ←	1	0	? ←
Nonstop car chases	0	0	1	1
Swords vs. karate	0	0	1	? ←

1  
0  
?



## Example applications

- 1. Did user  $j$  purchase an item after being shown? 1, 0, ?
- 2. Did user  $j$  fav/like an item? 1, 0, ?
- 3. Did user  $j$  spend at least 30sec with an item? 1, 0, ?
- 4. Did user  $j$  click on an item? 1, 0, ?

Meaning of ratings:

- 1 - engaged after being shown item
- 0 - did not engage after being shown item
- ? - item not yet shown

## From regression to binary classification

- Previously:
- Predict  $y^{(i,j)}$  as  $w^{(j)} \cdot x^{(i)} + b^{(j)}$
- For binary labels:  
Predict that the probability of  $y^{(i,j)} = 1$   
is given by  $g(w^{(j)} \cdot x^{(i)} + b^{(j)})$   
where  $g(z) = \frac{1}{1+e^{-z}}$

## Cost function for binary application

Previous cost function:

$$\frac{1}{2} \sum_{(i,j): r(i,j)=1} (\underbrace{w^{(j)} \cdot x^{(i)} + b^{(j)}}_{f(x)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (w_k^{(j)})^2$$

Loss for binary labels  $y^{(i,j)}$ :  $f_{(w,b,x)}(x) = g(w^{(j)} \cdot x^{(i)} + b^{(j)})$

$$L(f_{(w,b,x)}(x), y^{(i,j)}) = -y^{(i,j)} \log(f_{(w,b,x)}(x)) - (1 - y^{(i,j)}) \log(1 - f_{(w,b,x)}(x)) \quad \leftarrow \text{Loss for single example}$$

$$J(w, b, x) = \sum_{(i,j): r(i,j)=1} L(f_{(w,b,x)}(x), y^{(i,j)}) \quad \leftarrow \text{cost for all examples}$$

成功推广到二进制

# Mean normalization

归一化 for Run faster

## Users who have not rated any movies

Movie	Alice(1)	Bob (2)	Carol (3)	Dave (4)	Eve (5)
Love at last	5	5	0	0	?
Romance forever	5	?	?	0	?
Cute puppies of love	?	4	0	?	?
Nonstop car chases	0	0	5	4	?
Swords vs. karate	0	0	5	?	?

$$\begin{bmatrix} 5 & 5 & 0 & 0 & ? \\ 5 & ? & ? & 0 & ? \\ ? & 4 & 0 & ? & ? \\ 0 & 0 & 5 & 4 & ? \\ 0 & 0 & 5 & 0 & ? \end{bmatrix}$$

$$\min_{\substack{w^{(1)}, \dots, w^{(n_u)} \\ b^{(1)}, \dots, b^{(n_u)} \\ x^{(1)}, \dots, x^{(n_m)}}} \frac{1}{2} \sum_{(i,j): r(i,j)=1} (w^{(j)} \cdot x^{(i)} + b^{(j)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (w_k^{(j)})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2$$

为了进行均值归一化,  
To carry out mean normalization,

对新用户要用 mean norm... 因为不能全是0

## Mean Normalization

$$\begin{aligned} & \begin{bmatrix} 5 & 5 & 0 & 0 & ? \\ 5 & ? & ? & 0 & ? \\ ? & 4 & 0 & ? & ? \\ 0 & 0 & 5 & 4 & ? \\ 0 & 0 & 5 & 0 & ? \end{bmatrix} \begin{matrix} 2.5 \\ 2.5 \\ 2 \\ 2.25 \\ 1.25 \end{matrix} \\ & \mu = \begin{bmatrix} 2.5 \\ 2.5 \\ 2 \\ 2.25 \\ 1.25 \end{bmatrix} \quad \begin{bmatrix} 2.5 & 2.5 & -2.5 & -2.5 & ? \\ 2.5 & ? & ? & -2.5 & ? \\ ? & 2 & -2 & ? & ? \\ -2.25 & -2.25 & 2.75 & 1.75 & ? \\ -1.25 & -1.25 & 3.75 & -1.25 & ? \end{bmatrix} \end{aligned}$$

For user  $j$ , on movie  $i$  predict:

$$w^{(j)} \cdot x^{(i)} + b^{(j)} + \mu_i$$

$$y^{(i,j)} \downarrow \\ w^{(j)} \cdot b^{(j)} \cdot x^{(i)}$$

User 5 (Eve):

$$w^{(5)} = \begin{bmatrix} 0 \\ 0 \end{bmatrix} \quad b^{(5)} = 0 \quad \underbrace{w^{(5)} \cdot x^{(1)} + b^{(5)}}_{2.5} + \mu_1 = 2.5$$

2.5 rather than think Eve will rate all movie zero stars just because she

# TensorFlow implementation

## Derivatives in ML

### Gradient descent algorithm

Repeat until convergence


$$\underline{w} = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

$$\underline{b} = b - \alpha \frac{\partial}{\partial b} J(w, b) \leftarrow b = 0$$

Learning rate

Derivative

$J(w)$



$$f(x) = w \cdot x \quad b = 0$$

## Implementation in TensorFlow

Gradient descent algorithm

Repeat until convergence

$$\begin{aligned} w &= w - \alpha \frac{\partial}{\partial w} J(w, b, x) \\ b &= b - \alpha \frac{\partial}{\partial b} J(w, b, x) \\ x &= x - \alpha \frac{\partial}{\partial x} J(w, b, x) \end{aligned}$$

```
# Instantiate an optimizer.
optimizer = keras.optimizers.Adam(learning_rate=1e-1)

iterations = 200
for iter in range(iterations):
    # Use TensorFlow's GradientTape
    # to record the operations used to compute the cost
    # with tf.GradientTape() as tape:
        # Compute the cost (forward pass is included in cost)
        cost_value = cofiCostFuncV(X, W, b, Ynorm, R,
                                   num_users, num_movies, lambda)
        # Use the gradient tape to automatically retrieve
        # the gradients of the trainable variables with respect to
        # the loss
        grads = tape.gradient(cost_value, [X, W, b])
        # Run one step of gradient descent by updating
        # the value of the variables to minimize the loss.
        optimizer.apply_gradients(zip(grads, [X, W, b]))
```

Dataset credit: Harper and Konstan. 2015. The MovieLens Datasets: History and Context

## Finding related items

### Finding related items

The features  $x^{(i)}$  of item  $i$  are quite hard to interpret.

To find other items related to it,

find item  $k$  with  $x^{(k)}$  similar to  $x^{(i)}$

i.e. with smallest distance

$$\sum_{l=1}^n (x_l^{(k)} - x_l^{(i)})^2$$
$$\|x^{(k)} - x^{(i)}\|^2$$

romance

action

$x_1, x_2, x_3$   
 $n$

$x^{(k)}$   $x^{(i)}$



# Limitations of Collaborative Filtering

→ Cold start problem. How to

- • rank new items that few users have rated?
- • show something reasonable to new users who have rated few items?

→ Use side information about items or users:

- • Item: Genre, movie stars, studio, ....
- • User: Demographics (age, gender, location), expressed preferences, ... }

## Second Recommender System

### Collaborative Filtering Vs Content-based Filtering

#### Collaborative filtering vs Content-based filtering

→ Collaborative filtering:

Recommend items to you based on rating of users who gave similar ratings as you

→ Content-based filtering:

Recommend items to you based on features of user and item to find good match

$r(i, j) = 1$  if user  $j$  has rated item  $i$

$r(i, j)$  rating given by user  $j$  on item  $i$  (if defined)

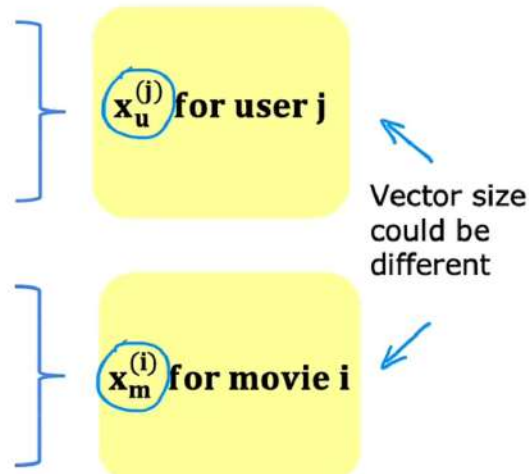
#### Examples of user and item features

##### User features:

- • Age
- • Gender (1 hot)
- • Country (1 hot, 200)
- • Movies watched (1000)
- • Average rating per genre
- ...

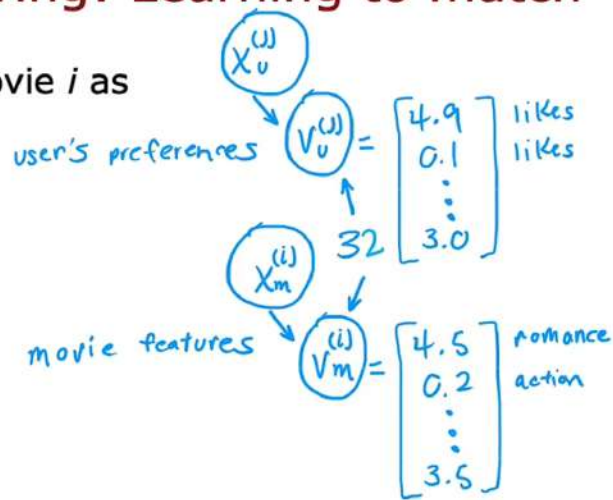

##### Movie features:

- • Year
- • Genre/Genres
- • Reviews
- • Average rating
- ...



# Content-based filtering: Learning to match


Predict rating of user  $j$  on movie  $i$  as




## Deep learning for content-based filtering

### Neural network architecture

$x_u \rightarrow v_u$  User network




$x_m \rightarrow v_m$  Movie network



Prediction:  $v_u^{(j)} \cdot v_m^{(i)}$   
 $g(v_u^{(j)} \cdot v_m^{(i)})$  to predict the probability that  $y^{(i,j)}$  is 1

### Neural network architecture



Cost function  $J = \sum_{(i,j):r(i,j)=1} (v_u^{(j)} \cdot v_m^{(i)} - y^{(i,j)})^2 + \text{NN regularization term}$



## Learned user and item vectors:

- $v_u^{(j)}$  is a vector of length 32 that describes user  $j$  with features  $x_u^{(j)}$
- $v_m^{(i)}$  is a vector of length 32 that describes movie  $i$  with features  $x_m^{(i)}$


To find movies similar to movie  $i$ :  $\|v_m^{(k)} - v_m^{(i)}\|^2$  small  
 $\|x^{(k)} - x^{(i)}\|^2$

Note: This can be pre-computed ahead of time

神经网络  
的神经

## Recommending from a large catalogue

How to efficiently find recommendation from a large set of items?



## Two steps: Retrieval & Ranking


Retrieval:

- • Generate large list of plausible item candidates ~100s
  - 1) For each of the last 10 movies watched by the user, find 10 most similar movies  
 $\|v_m^{(k)} - v_m^{(i)}\|^2$
  - 2) For most viewed 3 genres, find the top 10 movies
  - 3) Top 20 movies in the country
- • Combine retrieved items into list, removing duplicates and items already watched/purchased

## Two steps: Retrieval & ranking

### Ranking:

- Take list retrieved and rank using learned model



- Display ranked items to user

### Retrieval step

- • Retrieving more items results in better performance, but slower recommendations.
- • To analyse/optimize the trade-off, carry out offline experiments to see if retrieving additional items results in more relevant recommendations (i.e.,  $p(y^{(i,j)}) = 1$  of items displayed to user are higher).

100    500

## Ethical use of recommender systems

### What is the goal of the recommender system?

#### Recommend:


- • Movies most likely to be rated 5 stars by user
- • Products most likely to be purchased
- • Ads most likely to be clicked on *+ high bid*
- • Products generating the largest profit
- • Video leading to maximum watch time



# Ethical considerations with recommender systems


## Travel industry

Good travel experience  
to more users



## Payday loans

Squeeze customers  
more




Amelioration: Do not accept ads from exploitative businesses

## Other problematic cases:

- • Maximizing user engagement (e.g. watch time) has led to large social media/video sharing sites to amplify conspiracy theories and hate/toxicity
- Amelioration : Filter out problematic content such as hate speech, fraud, scams and violent content
- • Can a ranking system maximize your profit rather than users' welfare be presented in a transparent way?
- Amelioration : Be transparent with users

## Tensorflow implementation



```
user_NN = tf.keras.models.Sequential([  
    tf.keras.layers.Dense(256, activation='relu'),  
    tf.keras.layers.Dense(128, activation='relu'),  
    tf.keras.layers.Dense(32)  
])
```

```
item_NN = tf.keras.models.Sequential([  
    tf.keras.layers.Dense(256, activation='relu'),  
    tf.keras.layers.Dense(128, activation='relu'),  
    tf.keras.layers.Dense(32)  
])
```


```
# create the user input and point to the base network  
input_user = tf.keras.layers.Input(shape=(num_user_features))  
vu = user_NN(input_user)  
vu = tf.linalg.l2_normalize(vu, axis=1)
```

```
# create the item input and point to the base network  
input_item = tf.keras.layers.Input(shape=(num_item_features))  
vm = item_NN(input_item)  
vm = tf.linalg.l2_normalize(vm, axis=1)
```

```
# measure the similarity of the two vector outputs  
output = tf.keras.layers.Dot(axes=1)([vu, vm])
```

```
# specify the inputs and output of the model  
model = Model([input_user, input_item],
```

```
cost_fn = tf.keras.losses.MeanSquaredError)
```



也称为规范化向量的 l2

is also called normalizing the l2 norm of the vector,